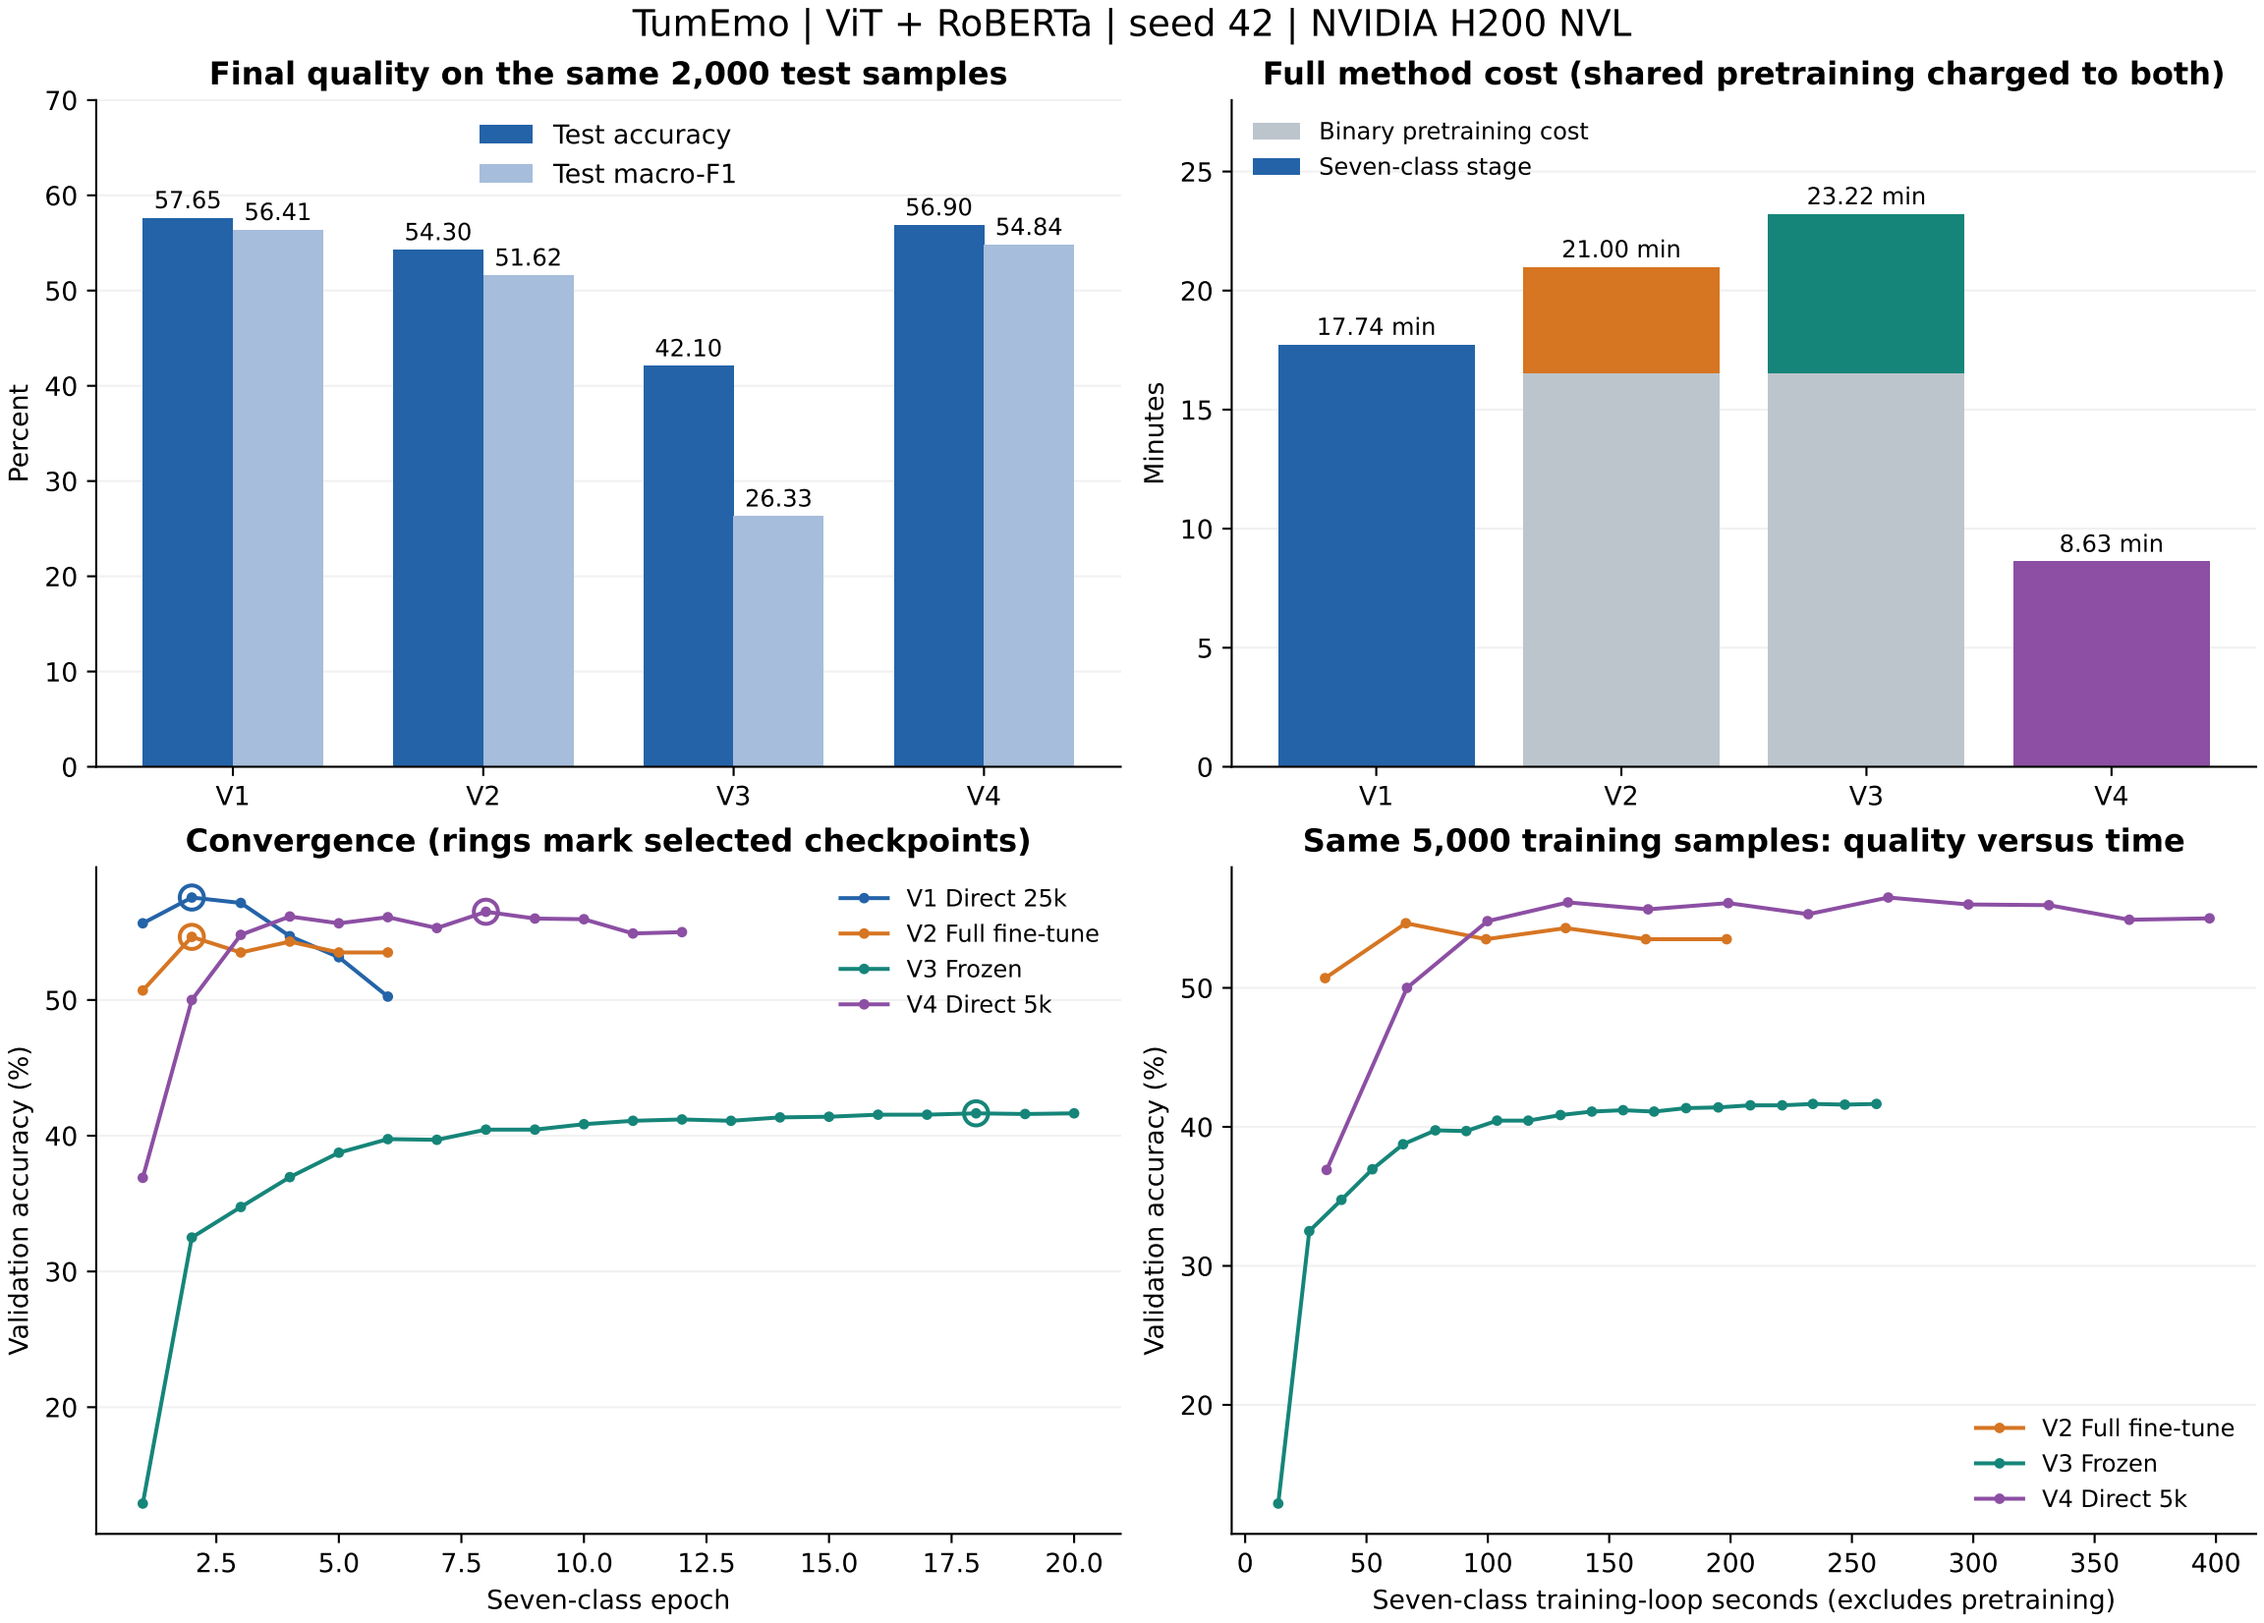
TumEmo | ViT + RoBERTa | seed 42 | NVIDIA H200 NVL
Final quality on the same 2,000 test samples
Test accuracy
Test macro-F1
Percent
57.65
56.41
54.30
51.62
42.10
26.33
56.90
54.84
V1
V2
V3
V4
Full method cost (shared pretraining charged to both)
Binary pretraining cost
Seven-class stage
Minutes
17.74 min
21.00 min
23.22 min
8.63 min
Convergence (rings mark selected checkpoints)
V1 Direct 25k
V2 Full fine-tune
V3 Frozen
V4 Direct 5k
Validation accuracy (%)
Seven-class epoch
Same 5,000 training samples: quality versus time
Seven-class training-loop seconds (excludes pretraining)